

相关章节：2.3.5

Robbins-Monro algorithm

已知 $p(z, \theta)$, 求 $f(\theta) = \mathbb{E}_{z|\theta}[z] = \int zp(z|\theta)dz = 0$ 的解 θ^* .

假设:

1. the conditional variance of z is finite so that $\mathbb{E}_{z|\theta}[(z - f)^2] < \infty$;
2. without loss of generality, consider the case where $f(\theta) > 0$ for $\theta > \theta^*$ and $f(\theta) < 0$ for $\theta < \theta^*$;

则由如下递推公式得到的序列 $\{\theta^{(N)}\}$ 依概率收敛于 θ^* :

$$\theta^{(N)} = \theta^{(N-1)} - a_{N-1}z(\theta^{(N-1)})$$

where $z(\theta^{(N)})$ is an observed value of z when θ takes the value $\theta^{(N)}$, 也就是说若记 $z_{N+1} = z(\theta^{(N)})$, 则 $z_{N+1} \sim p(z|\theta^N)$. The coefficients $\{a_N\}$ represent a sequence of positive numbers that satisfy the conditions:

$$\begin{aligned}\lim_{N \rightarrow \infty} a_N &= 0 \\ \sum_{N=1}^{\infty} a_N &= \infty \\ \sum_{N=1}^{\infty} a_N^2 &< \infty\end{aligned}$$

Note that the first condition $\lim_{N \rightarrow \infty} a_N = 0$ ensures that the successive corrections decrease in magnitude so that the process can converge to a limiting value. The second condition $\sum_{N=1}^{\infty} a_N = \infty$ is required to ensure that the algorithm does not converge short of the root, and the third condition $\sum_{N=1}^{\infty} a_N^2 < \infty$ is needed to ensure that the accumulated noise has finite variance and hence does not spoil convergence.

Sequential MLE

By definition, the maximum likelihood solution θ_{ML} is a stationary point of the log likelihood function and hence satisfies:

$$-\frac{\partial}{\partial \theta} \left\{ \frac{1}{N} \sum_{n=1}^N \ln p_{model}(x_n|\theta) \right\} \Big|_{\theta_{ML}} = 0$$

其中 $x_n \sim p_{real}(x)$, $p_{model}(x|\theta)$ 是人为设定的函数形式已知的概率分布, θ 为待定参数.

当 $N \rightarrow \infty$ 时, θ_{ML} 即为如下方程的解:

$$\mathbb{E}_{x \sim p_{real}(x)} \left[-\frac{\partial}{\partial \theta} \ln p_{model}(x|\theta) \right] = 0$$

记

$$z = g(x, \theta) = -\frac{\partial}{\partial \theta} \ln p_{\text{model}}(x|\theta)$$

则方程进一步可改写为：

$$\mathbb{E}_{z \sim p_{\text{real}}(z|\theta)}[z] = 0$$

其中, $p_{\text{real}}(z|\theta)$ 由如下过程给出：给定 θ , 采样 $x \sim p_{\text{real}}(x)$, 接着通过映射 $g(x, \theta)$ 即得到 z 。接下来, 即可应用 Robbins-Monro algorithm 求解上述方程：

$$\begin{aligned}\theta^{(N)} &= \theta^{(N-1)} - a_{N-1} z(\theta^{(N-1)}) \\ &= \theta^{(N-1)} - a_{N-1} z_N \\ &= \theta^{(N-1)} - a_{N-1} g(x_N, \theta^{(N-1)}) \\ &= \theta^{(N-1)} - a_{N-1} \frac{\partial}{\partial \theta^{(N-1)}} [-\ln p_{\text{model}}(x_N|\theta^{(N-1)})]\end{aligned}$$

注意, $x_N \sim p_{\text{real}}(x)$ 而不是 $x_N \sim p_{\text{model}}(x|\theta^{(N-1)})$, 也就是真实样本。上式即为 sequential MLE 的计算公式。

若 $p_{\text{model}}(x|\theta)$ 取高斯分布：

$$p_{\text{model}}(x|\theta) = p_{\text{model}}(x|\mu, \sigma) = \frac{1}{\sqrt{2\pi}\sigma} \exp \left\{ -\frac{(x - \mu)^2}{2\sigma^2} \right\}$$

这里我们只考察均值而假设方差已知, 可得：

$$z = g(x, \mu) = -\frac{x - \mu}{\sigma^2}$$

取 $a_{N-1} = \sigma^2/N$, 则：

$$\mu^{(N)} = \mu^{(N-1)} + \frac{1}{N} (x_N - \mu^{(N-1)})$$

其中, $x_N \sim p_{\text{real}}(x)$ 。可以看到, 上式与直接通过 MLE 得到的计算公式相同。

进一步, 若 $p_{\text{real}}(x)$ 真的是高斯分布：

$$p_{\text{real}}(x) = \mathcal{N}(x|\mu_*, \sigma_*^2)$$

又 $z = g(x, \mu) = -\frac{x - \mu}{\sigma^2}$, 已知 $p_{\text{real}}(z|\mu)$ 也服从高斯分布。此时, 方程 $\mathbb{E}_{z \sim p_{\text{real}}(z|\mu)}[z] = 0$ 可解析求解：

$$\begin{aligned}\mathbb{E}_{z \sim p_{\text{real}}(z|\mu)}[z] &= -\frac{\mu_* - \mu}{\sigma^2} = 0 \\ \Downarrow \\ \mu &= \mu_*\end{aligned}$$

又由 Robbins-Monro algorithm 可知上面得到的 $\{\mu^{(N)}\}$ 依概率收敛于方程 $\mathbb{E}_{z \sim p_{\text{real}}(z|\theta)}[z] = 0$ 的解, 即 μ_* 。这就表明, 若真实分布也为高斯分布, 我们求得的 $\{\mu^{(N)}\}$ 确为真实均值 μ_* 。